

Five Building Blocks of a Modern AI Agent

AGENTS.md, skills, MCP, A2A and subagents — what each one is for, and why they are not interchangeable.

For engineers building agents who need to know which mechanism solves which problem · 26 July 2026

The one- or two-page version: what matters, the topics it covers, and every term defined. Read the full lesson for the explanations and worked examples.

[Watch the source video](#)[Read the full lesson](#)[Read the condensed version](#)

What matters

- ✓ The model is not the agent — the instruction layer wrapped around it is what makes it one.
- ✓ AGENTS.md is a project-level instruction file, read whenever the agent starts work there.
- ✓ Nested AGENTS.md files stack, with those closer to the working directory taking precedence.
- ✓ A skill is a folder with a SKILL.md whose description tells the agent when to invoke it, so it stays out of context until needed.
- ✓ MCP is a standard interface between an agent and external tools or data, replacing per-target connectors.
- ✓ A2A is a standard for agents to discover and delegate to each other, using a published agent card.
- ✓ MCP is asymmetric — agent to tool. A2A is between peers, each autonomous.
- ✓ Subagents run in their own fresh context and return a result, keeping the parent's context clean and enabling parallel work.

What it covers

- 01 The model is not the agent** 0:00
A model generates text and reasons. What turns it into an agent is the instruction layer wrapped around it.
- 02 AGENTS.md: how this project works** 0:40
A Markdown file at the project root that the agent reads whenever it starts work there.
- 03 Nesting and precedence** 2:02
Several AGENTS.md files can coexist. The one closest to the working directory wins.
- 04 Agent skills: how to do a kind of task** 2:44
A folder with a SKILL.md whose description tells the agent when to pull it in — and which stays out of context until then.
- 05 MCP: reaching tools and data** 4:49
An open protocol that wraps any tool or data source in a standard interface, so one connector serves every agent.
- 06 A2A: agents talking to each other** 6:17
An open protocol for agent-to-agent communication, where each agent publishes a card describing what it does and how to reach it.
- 07 Subagents: work that will not fit one context** 8:24
A child agent spawned with its own fresh context, which does a piece of work and returns a result.
- 08 The five in one picture** 10:28
Two shape behaviour, two reach outward, one scales past a single context — and the symptoms tell you which is missing.

Glossary

AGENTS.md — A Markdown file at a project root holding commands, conventions and prohibitions, read whenever an agent starts work in that project.

Nesting / precedence — Several AGENTS.md files applied from the root down, with the file closest to the working directory overriding those above it.

SKILL.md — The file defining a skill: metadata including the description that tells the agent when to invoke it, and the steps to follow.

MCP — Model Context Protocol: an open protocol connecting AI applications to tools, data sources and workflows through a standard interface.

A2A — Agent to agent: an open protocol letting autonomous agents discover one another and delegate work.

Subagent — A child agent spawned by a parent, running in its own fresh context, doing one piece of work and returning a result.

Embarrassingly parallel — Work that splits into independent pieces needing no coordination — the case where fanning out to subagents pays.

Agentic AI Foundation — The body under the Linux Foundation that now governs several of these standards, including AGENTS.md and MCP.

Agent skill — A folder containing a SKILL.md plus scripts and resources, describing how to perform one kind of task.

Skill description — The one visible part of a skill before loading. It functions as a routing key, so a vague description makes the skill unreachable.

MCP server — The component wrapping a tool or data source so any MCP-speaking agent can use it, and the right place to enforce permission.

Agent card — A published description of what an agent does and how to talk to it, which other agents read in order to delegate.

Context isolation — The property that makes subagents useful: the child holds the bulk of the material and the parent receives only the conclusion.

Explanations are AI-generated. Check anything you intend to rely on.