

Five AI Risks That Can Get Someone Fired

Recognize shadow AI, data leakage, hallucination laundering, prompt injection, and unauthorized agents before they turn into a career-ending incident.

For employees and IT leaders who use or deploy AI tools at work and want to avoid preventable, accountability-shifting mistakes · 27 July 2026

The one- or two-page version: what matters, the topics it covers, and every term defined. Read the full lesson for the explanations and worked examples.

[Watch the source video](#)[Read the full lesson](#)[Read the condensed version](#)

What matters

- ✓ Shadow AI is any AI tool used for work that IT never vetted or approved; per IBM's Cost of a Data Breach report, 1 in 5 organizations have reported a breach caused by it.
- ✓ Data pasted into an unapproved AI tool can be used to train a future model version, at which point it cannot practically be retrieved or deleted.
- ✓ Banning AI tools outright tends to push usage onto personal devices or newer unblocked tools, so IT loses visibility instead of gaining control.
- ✓ Hallucination laundering is copying an AI's confidently wrong output into a real work product without checking it — the mistake becomes the submitter's, not the AI's.
- ✓ Direct prompt injection is typed straight into a chat interface; indirect prompt injection hides malicious instructions inside a document, email, or web page the AI later retrieves.
- ✓ Indirect prompt injection is generally more dangerous because there is no suspicious user input to flag — the attack rides in with otherwise legitimate data.
- ✓ A zombie AI agent is one left running and authenticated after its project ended; it becomes an unmonitored backdoor even when it was built with no bad intent.
- ✓ A written AI governance policy — naming approved tools, off-limits data, and required human review — closes most of these gaps without banning AI outright.

What it covers

01 Shadow AI: Unapproved Tools Inside the Company 0:00

Shadow AI is any AI tool employees use for work that corporate IT has neither vetted nor approved, and it is already causing real data breaches.

02 Data Leakage: Where the Pasted Content Actually Goes 1:27

Content pasted into an unapproved AI tool is sent to a third-party server and can end up used to train a future model, at which point it can't be clawed back.

03 AI Governance: The Actual Fix for Shadow AI and Leakage 2:48

A written policy naming which AI tools are approved, how they may be used, and what data is off-limits closes the shadow AI gap without banning AI outright.

04 Hallucination Laundering: Presenting Fabricated Output as Fact 3:31

Hallucination laundering happens when someone copies an AI's confidently wrong output straight into their own work without checking it, turning a model's mistake into their personal, career-attached mistake.

05 Prompt Injection: Direct and Indirect Attacks 4:55

Prompt injection is an attack that gets an AI system to ignore its original instructions, either by typing the attack directly into a chat box or by hiding it inside content the AI is asked to read.

06 Unauthorized Agentic AI and the Zombie Agent Problem 7:04

An AI agent that can act autonomously — reading databases, calling APIs, sending messages — becomes dangerous when employees spin it up without oversight and it keeps running, authenticated, long after anyone remembers it exists.

Glossary

Shadow AI — Any AI tool used for work that corporate IT has not vetted or approved, such as a personal chatbot account or an unsanctioned browser plugin.

AI governance — A written policy naming which AI tools are approved, how they may be used, and which data categories may never be submitted to them.

Hallucination laundering — Copying unverified, hallucinated AI output into a real work product and submitting it as one's own verified work.

Direct prompt injection — A malicious instruction typed straight into an AI system's chat interface by the attacker.

System prompt — The set of ground-rule instructions given to an AI model before any user input, defining what it should and shouldn't do.

Zombie AI agent — An AI agent left running and authenticated after the project that created it has ended, becoming an unmonitored backdoor into internal systems.

Data leakage — Sensitive content sent to a third-party AI tool outside the organization's control, where it may be retained or used for training.

Hallucination — AI-generated content that sounds plausible and is delivered confidently but is factually incorrect.

Prompt injection — An attack that crafts input to override an AI system's original instructions, such as a chatbot's system prompt.

Indirect prompt injection — A malicious instruction hidden inside content the AI retrieves and processes, such as a document, email, or web page, rather than typed by the attacker directly.

Agentic AI — AI that works autonomously toward a goal, capable of actions like reading and writing data, calling APIs, executing code, and sending messages without step-by-step human approval.

Explanations are AI-generated. Check anything you intend to rely on.