

LLM-as-a-Judge: Direct Assessment, Pairwise Comparison & the Biases That Skew Both

Choose the right judging strategy for your task, and recognise the three biases that make a judge quietly wrong.

For engineers and researchers who need to evaluate large volumes of generated text without labelling every output by hand · 27 July 2026

The one- or two-page version: what matters, the topics it covers, and every term defined. Read the full lesson for the explanations and worked examples.

[Watch the source video](#)

[Read the full lesson](#)

[Read the condensed version](#)

What matters

- ✓ LLM-as-a-judge enables reference-free evaluation: judging quality without a gold-standard answer to compare against.
- ✓ Direct assessment scores each output independently against a rubric you design; pairwise comparison asks which of two outputs is better.
- ✓ With multiple candidates, pairwise judgements can be combined by a ranking algorithm to produce an overall ordering.
- ✓ In IBM's user research on the open-source EvalAssist framework, about half preferred direct assessment for its clarity and control, a quarter preferred pairwise for subjective tasks, and the rest combined both.
- ✓ The three advantages over traditional metrics are scale, flexibility as your criteria evolve, and the ability to judge nuanced qualities like whether text sounds natural.
- ✓ Positional bias: the judge favours an output because of where it appeared, not because it is better.
- ✓ Verbosity bias: the judge treats length as a proxy for quality and prefers longer answers even when they ramble.
- ✓ Self-enhancement bias: the judge favours output it recognises as its own.
- ✓ Positional swaps — running the same comparison with the order reversed — detect positional bias directly. Bias means stay vigilant, not that the system is broken.

What it covers

- 01 Why traditional metrics run out** 0:00
Manual labelling does not scale and word-overlap metrics need a reference answer they cannot always have.
- 02 Strategy one: direct assessment** 0:42
Design a rubric and score each output against it independently, one at a time.
- 03 Strategy two: pairwise comparison** 1:23
Instead of a standalone label, ask which of two outputs is better — then rank many outputs from the accumulated comparisons.
- 04 Which strategy do people actually choose?** 2:06
User research on the open-source EvalAssist framework found the split was roughly half direct assessment, a quarter pairwise, and the rest a combination — with the choice depending on both task and user.
- 05 Why use an LLM judge: scale, flexibility, nuance** 2:47

It handles volume, it adapts as your criteria shift, and it can evaluate qualities no overlap metric reaches.

06 **The three biases that skew a judge** 3:29

Positional, verbosity, and self-enhancement bias each make a judge favour outputs for reasons unrelated to quality.

07 **Catching bias with positional swaps** 4:54

Run the same comparison twice with the options in reverse order; if the verdict flips, the judgement was driven by position rather than content.

Glossary

LLM-as-a-judge — Using a language model to evaluate the outputs of another model, applying a criterion you specify.

Direct assessment — A judging strategy where each output is scored independently against a rubric you designed.

Rubric — The criterion being evaluated plus the set of options the judge may return.

BLEU / ROUGE — Traditional text metrics that score an output by how much its wording overlaps with a reference answer.

Verbosity bias — A judge favouring longer outputs regardless of quality, treating length as a signal of thoroughness.

Positional swap — Re-running a pairwise comparison with the two options in reverse order to test whether the verdict was driven by position.

Reference-free evaluation — Judging an output on its own merits, without needing a gold-standard answer to compare it against.

Pairwise comparison — A judging strategy where two outputs are shown together and the judge picks which is better.

Ranking algorithm — A method for converting many individual pairwise verdicts into an overall ordering of all candidates.

Positional bias — A judge favouring an output because of the position it appeared in rather than its content.

Self-enhancement bias — A judge rating its own generated output more highly than others', independent of actual quality.

EvalAssist — IBM's open-sourced evaluation framework, used in the user research on how people choose between judging strategies.

Explanations are AI-generated. Check anything you intend to rely on.